

Basic NLP Text Processing

MIE223
Winter 2025

1 Basic NLP Text Processing. NLP=Natural Language Processing

1.1 NLP Text Processing Pipeline

nlTK provides implementations for most operations

- Document → Sections and Paragraphs
- Paragraphs → Sentences (sentence segmentation / extraction)
- Sentences → Tokens
- Tokens → Lemmas or Morphological Variants / Stems
- Tokens → Part-of-speech (POS) Tags
- Tokens, POS Tags → Phrase Chunks (Noun & Verb Phrases)
- Tokens, POS Tags → Parse Trees
 - Augment above with coreference, entailment, sentiment, ...

2 Basic Text Processing: Word tokenization

2.1 Text Normalization

Every NLP task needs to do text normalization:

1. Segmenting/tokenizing words in running text
2. Normalizing word formats
3. Segmenting sentences in running text

2.2 How many words?

I do uh main- mainly business data processing

- Fragments, filled pauses

Seuss's cat in the hat is different from other cats

- Lemma: same stem, part of speech, rough word sense
- e.g. cat and cats = same lemma
- Wordform: the full inflected surface form
- e.g. cat and cats = different wordforms

they lay back on the San Francisco grass and looked at the stars and their

- Type: an element of the vocabulary.
- Token: an instance of that type in running text.
- How many? Depends on how you segment. San Francisco can be one token

- 15 tokens (or 14)
- 13 types (or 12) (or 11?)

N = number of tokens (1)

V = vocabulary = set of types (2)

$|V|$ = size of vocabulary (3)

$|V| > O(N^{1/2})$ (4)

	Tokens = N	Types = $ V $
Switchboard phone conversations	2.4 million	20 thousand
Shakespeare	884,000	31 thousand
Google N-grams	1 trillion	13 million

2.3 Issues in Tokenization

- Finland's capital → Finland Finlands Finland's ?
- what're, I'm, isn't → What are, I am, is not
- Hewlett-Packard → Hewlett Packard ?
- state-of-the-art → state of the art ?
- Lowercase → lower-case lowercase lower case ?
- San Francisco → **one token or two?**
- m.p.h., PhD. → ??


Keep hyphens together? Remove contractions? Does upper and lower case matter?

2.4 Tokenization: language issues

- French
 - L'ensemble → one token or two?
 - L ? L' ? Le ?
 - Want l'ensemble to match with un ensemble
- German noun compounds are not segmented
 - Lebensversicherungsgesellschaftsangestellter
 - 'life insurance company employee'
 - German information retrieval needs compound splitter



- Chinese and Japanese no spaces between words:
 - 莎拉波娃现在居住在美国东南部的佛罗里达。
 - 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
 - Sharapova now lives in US southeastern Florida
- Further complicated in Japanese, with multiple alphabets intermingled
 - Dates/amounts in multiple formats



10 End-user can express query entirely in hiragana!

3 Word Normalization and Stemming

3.1 Normalization

- Need to “normalize” terms
 - Information Retrieval: indexed text & query terms must have same form.
 - * We want to match U.S.A. and USA
- We implicitly define equivalence classes of terms
 - e.g., deleting periods in a term
- Alternative: asymmetric expansion:
 - Enter: window Search: window, windows
 - Enter: windows Search: Windows, windows, window
 - Enter: Windows Search: Windows
- Potentially more powerful, but less efficient

3.2 Case folding

- Applications like IR: reduce all letters to lower case
 - Since users tend to use lower case
 - Possible exception: upper case in mid-sentence?
 - * e.g., General Motors
 - * Fed vs. fed
 - * SAIL vs. sail
- For sentiment analysis, MT, Information extraction
 - Case is helpful (US versus us is important)

3.3 Lemmatization

- Reduce inflections or variant forms to base form
 - am, are, is → be
 - car, cars, car's, cars' → car
- the boy's cars are different colors → the boy car be different color
- Lemmatization: have to find correct dictionary headword form
- Machine translation
 - Spanish quiero ('I want'), quieres ('you want') same lemma as querer 'want'

3.4 Morphology

- Morphemes:
 - The small meaningful units that make up words
 - Stems: The core meaning-bearing units
 - Affixes: Bits and pieces that adhere to stems
 - * Often with grammatical functions

3.5 Stemming

- Reduce terms to their stems in information retrieval
- Stemming is crude chopping of affixes
 - language dependent
 - e.g., automate(s), automatic, automation all reduced to automat.
- for example compressed and compression are both accepted as equivalent to compress. → for example compress and compress are both accepted as equivalent to compress

3.6 Porter's algorithm: The most common English stemmer

Step 1a

sses → ss	caresses → caress
ies → i	ponies → poni
ss → ss	caress → caress
s → ∅	cats → cat

Step 1b

(*v*)ing → ∅	walking → walk
	sing → sing
(*v*)ed → ∅	plastered → plaster
...	

Step 2 (for long stems)

ational → ate	relational → relate
izer → ize	digitizer → digitize
ator → ate	operator → operate
...	

Step 3 (for longer stems)

al → ∅	revival → reviv
able → ∅	adjustable → adjust
ate → ∅	activate → activ
...	

3.7 Viewing morphology in a corpus Why only strip –ing if there is a vowel?

(*v*)ing → ∅ walking → walk sing → sing

(*v*)ing → ∅ walking → walk
sing → sing

```
tr -sc 'A-Za-z' '\n' < shakes.txt | grep 'ing$' | sort | uniq -c | sort -nr
```

1312 King	548 being
548 being	541 nothing
541 nothing	152 something
388 king	145 coming
375 bring	130 morning
358 thing	122 having
307 ring	120 living
152 something	117 loving
145 coming	116 Being
130 morning	102 going

```
tr -sc 'A-Za-z' '\n' < shakes.txt | grep '[aeiou].*ing$' | sort | uniq -c | sort -nr
```

3.8 Dealing with complex morphology is sometimes necessary

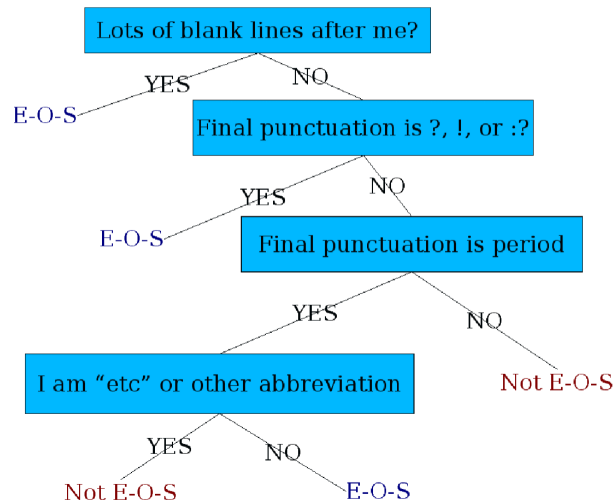
- Some languages requires complex morpheme segmentation
- Turkish
- Uygarlastiramadiklarimizdanmissinizcasina
- ‘(behaving) as if you are among those whom we could not civilize’
 - Uygar ‘civilized’ + las ‘become’
 - tir ‘cause’ + ama ‘not able’
 - dik ‘past’ + lar ‘plural’
 - imiz ‘p1pl’ + dan ‘abl’
 - mis ‘past’ + siniz ‘2pl’ + casina ‘as if’

4 Sentence Segmentation and Decision Trees

4.1 Sentence Segmentation

- !, ? are relatively unambiguous
- Period “.” is quite ambiguous
 - Sentence boundary
 - Abbreviations like Inc. or Dr.
 - Numbers like .02% or 4.3
- Build a binary classifier
- f(input) = True, False
- ”The car is traveling 10 m.p.h.”
- Check which period is ending the sentence
 - Looks at a “.”
 - Decides EndOfSentence/NotEndOfSentence
 - Classifiers: hand-written rules, regular expressions, or machine-learning

4.2 Determining if a word is end-of-sentence: a Decision Tree



4.3 More sophisticated decision tree features

- Case of word with “.”: Upper, Lower, Cap, Number
- Case of word after “.”: Upper, Lower, Cap, Number
- Numeric features
 - Length of word with “.”
 - Probability(word with “.” occurs at end-of-s)
 - Probability(word after “.” occurs at beginning-of-s)

5 Regular Expressions: Detecting word pattern variations

5.1 Regular expressions

- A formal language for specifying text strings
- How can we search for any of these?
 - woodchuck
 - woodchucks
 - Woodchuck
 - Woodchucks

5.2 Regular Expressions: Disjunctions

Letters inside square brackets []

Pattern	Matches
[wW]oodchuck	Woodchuck, woodchuck
[1234567890]	Any digit

Ranges [A-Z]

Pattern	Matches	
[A-Z]	An upper case letter	<u>D</u> renched Blossoms
[a-z]	A lower case letter	<u>m</u> y beans were impatient
[0-9]	A single digit	Chapter <u>1</u> : Down the Rabbit Hole

5.3 Regular Expressions: Negation in Disjunction

Negations [^s]

Carat means negation only when first in []

Pattern	Matches	
[^A-Z]	Not an upper case letter	O <u>y</u> fn pripetchik
[^Ss]	Neither 'S' nor 's'	<u>I</u> have no exquisite reason"
[^e^]	Neither e nor ^	<u>L</u> ook here
a^b	The pattern a carat b	Look up <u>a^b</u> now

5.4 Regular Expressions: More Disjunction

Woodchucks is another name for groundhog! The pipe — for disjunction.

Pattern	Alternative
groundhog woodchuck	
yours mine	
a b c	= [abc]
[gG]roundhog [Ww]oodchuck	

5.5 Regular Expressions: ? * + .

Pattern	Matches	
colou?r	Optional previous char	<u>color</u> <u>colour</u>
oo*h!	0 or more of previous char	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
o+h!	1 or more of previous char	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
baa+		<u>baa</u> <u>baaa</u> <u>baaaa</u> <u>baaaaa</u>
beg.n		<u>begin</u> <u>begun</u> <u>begun</u> <u>beg3n</u>



Stephen C Kleene

Kleene *, Kleene +

5.6 Regular Expressions: Anchors: ^ \$

Pattern	Matches
^[A-Z]	<u>P</u> alo Alto
^[^A-Za-z]	<u>1</u> <u>"Hello"</u>
\. \$	The end <u>.</u>
. \$	The end <u>?</u> The end <u>!</u>

5.7 Example

Find me all instances of the word "the" in a text.

the: Misses capitalized examples

[tT]he: Incorrectly returns other or theology

[^a-zA-Z][tT]he[^a-zA-Z]

5.8 Errors

- The process we just went through was based on fixing two kinds of errors
 - Matching strings that we should not have matched (there, then, other)
 - * False positives (Type I)
 - Not matching things that we should have matched (The)
 - * False negatives (Type II)

- In NLP we are always dealing with these kinds of errors.
- Reducing the error rate for an application often involves two antagonistic efforts:
 - Increasing accuracy or precision (minimizing false positives)
 - Increasing coverage or recall (minimizing false negatives).

5.9 Exercise

- Write a regular expression to match dates
 - November 9, 1989
 - 17 December 1967
 - 11-09-1989 (likely this form on midterm)
 - 12/17/67
- Write a regular expression to match time expressions
 - Next Wednesday at noon
 - Tomorrow morning
 - Can't really as there's no general pattern

5.10 Regex Summary

- Regular expressions play a surprisingly large role
 - Sophisticated sequences of regular expressions are often the first model for any text processing text
- For many hard tasks, we use machine learning classifiers
 - But regular expressions are used as features in the classifiers
 - Can be very useful in capturing generalizations

5.11 On the exam

- Match dates from 2000 to 2025 (Reg-Ex)
 - Only compound statements allowed
 - Do not just list out all the options
 - 20[0-1][0-9] — 202[0-5]