

Introduction to Privacy and Anonymity

MIE223
Winter 2025

1 Privacy and Anonymity

1.1 AOL Privacy Debacle

- In August 2006, AOL released anonymized search query logs
 - 657K users, 20M queries over 3 months (March-May)
- Opposing goals
 - Analyze data for research purposes, provide better services for users and advertisers
 - Protect privacy of AOL users
 - * Government laws and regulations
 - * Search queries may reveal income, evaluations, intentions to acquire goods and services, etc.

1.2 AOL User 4417749

- AOL query logs have the form $\langle \text{AnonID}, \text{Query}, \text{QueryTime}, \text{ItemRank}, \text{ClickURL} \rangle$
 - ClickURL is the truncated URL
- NY Times re-identified AnonID 4417749
 - Sample queries: “numb fingers”, “60 single men”, “dog that urinates on everything”, “landscapers in Lilburn, GA”, several people with the last name Arnold
 - * Lilburn area has only 14 citizens with the last name Arnold
 - NYT contacts the 14 citizens, finds out AOL User 4417749 is 62-year-old Thelma Arnold

1.3 Foundations of Privacy

- Consent:
 - GDPR (EU), US (Privacy Act of 1974)
- Notice: you have to accept collection practices
 - Question: who are some of the major providers of user web data?
- De-identification
 - Only release attributes that could not identify you
 - Historically founded on principle of k-anonymity
 - * Re-identification: multiple attributes can as well as quasi- identifiers (partial postal code) that link you accross datasets (medical, voter) even with k-anonymity
 - Sweeney (Harvard) used 1990 Census data to estimate that 0.04 percent of the United States population was uniquely identified by the basic demographic fields allowed by the HIPAA Safe Harbor – namely, year of birth, gender, and first 3 digits of ZIP

1.4 Background

- Large amount of person-specific data has been collected in recent years
 - Both by governments and by private entities
- Data and knowledge extracted by data mining techniques represent a key asset to the society
 - Analyzing trends and patterns
 - Formulating public policies
- Laws and regulations require that some collected data must be made public
 - For example, Census data

1.5 Public Data Conundrum

- Health-care datasets
 - Clinical studies, hospital discharge databases
- Genetic datasets
 - \$1000 genome, HapMap, deCode
- Demographic datasets
 - U.S. Census Bureau, sociology studies
- Search logs, recommender systems, social networks, blogs
 - AOL search data, social networks of blogging sites, Netflix movie ratings, Amazon

1.6 What About Privacy?

- First thought: anonymize the data
- How?
- Remove “personally identifying information” (PII)
 - Name, Social Security number, phone number, email, address... what else?
 - Anything that identifies the person directly
- Is this enough? No!

1.7 Re-identification by Linking

Microdata					Voter registration data			
ID	QID	SA			Name	Zipcode	Age	Sex
Name	Zipcode	Age	Sex	Disease				
Alice	47677	29	F	Ovarian Cancer	Alice	47677	29	F
Betty	47602	22	F	Ovarian Cancer	Bob	47983	65	M
Charles	47678	27	M	Prostate Cancer	Carol	47677	22	F
David	47905	43	M	Flu	Dan	47532	23	M
Emily	47909	52	F	Heart Disease	Ellen	46789	43	F
Fred	47906	47	M	Heart Disease				

ID is the PII, the rest isn't. When combined they are quasi-identifiers
Voter registration data is all PII.

1.8 Latanya Sweeney's Attack (1997)

Massachusetts hospital discharge dataset

Medical Data Released as Anonymous

SSN	Name	City	Date Of Birth	Sex	ZIP	Marital Status	Problem
			09/27/64	female	02139	divorced	hypertension
			09/30/64	female	02139	divorced	obesity
	asian		04/15/64	male	02139	married	chest pain
	asian		04/15/64	male	02139	married	obesity
	black		03/13/63	male	02138	married	hypertension
	black		03/18/63	male	02138	married	shortness of breath
	black		09/13/64	female	02141	married	shortness of breath
	black		09/07/64	female	02141	married	obesity
	white		05/14/61	male	02138	single	chest pain
	white		05/08/61	male	02138	single	obesity
	white		09/15/61	female	02142	widow	shortness of breath

Voter List

Name	Address	City	ZIP	DOB	Sex	Party
Sue J. Carlson	1459 Main St.	Cambridge	02142	9/15/61	female	democrat

Figure 1.8- Identifying anonymous data by linking to external data

Public voter dataset

1.9 Quasi-Identifiers

- Key attributes
 - Name, address, phone number - uniquely identifying!
 - Always removed before release
- Quasi-identifiers
 - (5-digit ZIP code, birth date, gender) uniquely identify 87% of the population in the U.S.
 - Can be used for linking anonymized dataset with other datasets

1.10 Classification of Attributes

- Sensitive attributes
 - Medical records, salaries, etc.
 - These attributes are what the researchers need, so they are always released directly


Key Attribute (often PII)	Quasi-identifier			Sensitive attribute
Name	DOB	Gender	Zipcode	Disease
Andre	1/21/76	Male	53715	Heart Disease
Beth	4/13/86	Female	53715	Hepatitis
Carol	2/28/76	Male	53703	Brochitis
Dan	1/21/76	Male	53703	Broken Arm
Ellen	4/13/86	Female	53706	Flu
Eric	2/28/76	Female	53706	Hang Nail

1.11 K-Anonymity: Intuition

- The information for each person contained in the released table cannot be distinguished from at least k-1 individuals whose information also appears in the release
 - Example: you try to identify a man in the released table, but the only information you have is his birth date and gender. There are k men in the table with the same birth date and gender.
- Any quasi-identifier present in the released table must appear in at least k records

1.12 k-Anonymity via Generalization

- Goal of k-Anonymity
 - Each record is indistinguishable from at least k-1 other records
 - These k records form an equivalence class
- Generalization: replace quasi-identifiers with less specific, but semantically consistent values
- If you just drop rows if not k-anonymous you cause missingness not at random
- you want to avoid dropping rows to avoid changing decisions, you can just generate attributes instead
- Example: replace 5-digit ZIP code with 3-digit ZIP code



1.13 Achieving k-Anonymity

- Generalization
 - Replace specific quasi-identifiers with less specific values until get k identical values
 - Partition ordered-value domains into intervals
- Problem: Suppression
 - When generalization causes too much information loss
 - This is common with “outliers”
- Lots of algorithms in the literature
 - Aim to produce “useful” anonymizations
 - ... usually without any clear notion of utility

1.14 Example of a k-Anonymous Table

	Race	Birth	Gender	ZIP	Problem
t1	Black	1965	m	0214*	short breath
t2	Black	1965	m	0214*	chest pain
t3	Black	1965	f	0213*	hypertension
t4	Black	1965	f	0213*	hypertension
t5	Black	1964	f	0213*	obesity
t6	Black	1964	f	0213*	chest pain
t7	White	1964	m	0213*	chest pain
t8	White	1964	m	0213*	obesity
t9	White	1964	m	0213*	short breath
t10	White	1967	m	0213*	chest pain
t11	White	1967	m	0213*	chest pain

Figure 2 Example of k-anonymity, where $k=2$ and $QI=\{Race, Birth, Gender, ZIP\}$

$k = 2$ and not $k = 3$ because we take the lowest k value.

1.15 Example of Generalization (1)

Race	Birth	Gender	ZIP	Problem
t1 Black	1965	m	0214*	short breath
t2 Black	1965	m	0214*	chest pain
t3 Black	1965	f	0213*	Hypertension
t4 Black	1965	f	0213*	Hypertension
t5 Black	1964	f	0213*	obesity
t6 Black	1964	f	0213*	chest pain
t7 White	1964	m	0213*	chest pain
t8 White	1964	m	0213*	obesity
t9 White	1964	m	0213*	short breath
t10 White	1967	m	0213*	chest pain
t11 White	1967	m	0213*	chest pain

Name	Birth	Gender	ZIP	Race
Andre	1964	m	02135	White
Beth	1964	f	55410	Black
Carol	1964	f	90210	White
Dan	1967	m	02174	White
Ellen	1968	f	02237	White

By linking these 2 tables, you still don't learn Andre's problem

1.16 Example of Generalization (2)

Microdata

QID	SA		
Zipcode	Age	Sex	Disease
47677	29	F	Ovarian Cancer
47602	22	F	Ovarian Cancer
47678	27	M	Prostate Cancer
47905	43	M	Flu
47909	52	F	Heart Disease
47906	47	M	Heart Disease

Generalized table

QID	SA		
Zipcode	Age	Sex	Disease
476**	2*	*	Ovarian Cancer
476**	2*	*	Ovarian Cancer
476**	2*	*	Prostate Cancer
4790*	[43,52]	*	Flu
4790*	[43,52]	*	Heart Disease
4790*	[43,52]	*	Heart Disease

Released table is 3-anonymous

Technically yes, but what does k-Anonymity miss?

If the adversary knows Alice's quasi-identifier (47677, 29, F), they still do not know which of the first 3 records corresponds to Alice's record

1.17 Curse of Dimensionality

- Generalization fundamentally relies on locality of quasi-identifiers
 - Each record must have k close neighbors
- Real-world datasets are very sparse
 - Many attributes (dimensions)
 - Netflix Prize dataset: 17,000 dimensions
 - Amazon customer records: several million dimensions
 - "Nearest neighbor" is very far
- Projection to low dimensions loses all info
 - k-anonymized datasets are useless

1.18 HIPAA Privacy Rule (US)

"Under the safe harbor method, covered entities must remove all of a list of 18 enumerated identifiers and have no actual knowledge that the information remaining could be used, alone or in combination, to identify a subject of the information."

"The identifiers that must be removed include direct identifiers, such as name, street address, social security number, as well as other identifiers, such as birth date, admission and discharge dates, and five-digit zip code. The safe harbor requires removal of geographic subdivisions smaller than a State, except for the initial three

digits of a zip code if the geographic unit formed by combining all zip codes with the same initial three digits contains more than 20,000 people. In addition, age, if less than 90, gender, ethnicity, and other demographic information not listed may remain in the information. The safe harbor is intended to provide covered entities with a simple, definitive method that does not require much judgment by the covered entity to determine if the information is adequately de-identified.”

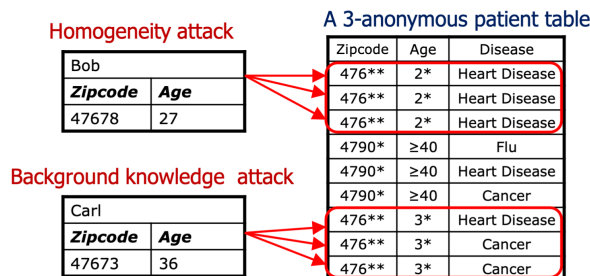
1.19 Two (and a Half) Interpretations

1. Membership disclosure: Attacker cannot tell that a given person in the dataset
2. Sensitive attribute disclosure: Attacker cannot tell that a given person has a certain sensitive attribute
3. Identity disclosure: Attacker cannot tell which record corresponds to a given person

This (3) interpretation is correct, assuming the attacker does not know anything other than quasi-identifiers. But this does not imply any privacy! Example: k clinical records, all HIV+

1.20 Attacks on k-Anonymity

- k-Anonymity does not provide privacy if
 - Sensitive values in an equivalence class lack diversity
 - The attacker has background knowledge



1.21 l-Diversity

Caucas	787XX	Flu
Caucas	787XX	Shingles
Caucas	787XX	Acne
Caucas	787XX	Flu
Caucas	787XX	Acne
Caucas	787XX	Flu
Asian/African American	78XXX	Flu
Asian/African American	78XXX	Flu
Asian/African American	78XXX	Acne
Asian/African American	78XXX	Shingles
Asian/African American	78XXX	Acne
Asian/African American	78XXX	Flu

Sensitive attributes must be "diverse" within each quasi-identifier equivalence class

1.22 Distinct l-Diversity

- Each equivalence class of quasi-identifiers has at least l well-represented sensitive values
- Doesn't prevent probabilistic inference attacks

...	Disease
...	...
...	HIV
...	HIV
...	...
...	HIV
...	pneumonia
...	bronchitis
...	...

10 records { 8 records have HIV
2 records have other values

1.23 Other Versions of l-Diversity

- Probabilistic l-diversity
 - The frequency of the most frequent value in an equivalence class is bounded by $1/l$
- Entropy l-diversity
 - The entropy of the distribution of sensitive values in each equivalence class is at least $\log(l)$
- Recursive (c,l)-diversity
 - $r_1 < c(r_l + r_{l+1} + \dots + r_m)$, where r_i is the frequency of the i -th most frequent value
 - Intuition: the most frequent value does not appear too frequently

1.24 t-Closeness

Distribution of sensitive attributes within each quasi-identifier group should be “close” to their distribution in the entire original database

Caucas	787XX	Flu
Caucas	787XX	Shingles
Caucas	787XX	Acne
Caucas	787XX	Flu
Caucas	787XX	Acne
Caucas	787XX	Flu
Asian/AfrAm	78XXX	Flu
Asian/AfrAm	78XXX	Flu
Asian/AfrAm	78XXX	Acne
Asian/AfrAm	78XXX	Shingles
Asian/AfrAm	78XXX	Acne
Asian/AfrAm	78XXX	Flu

1.25 Anonymous, “t-Close” Dataset

This is k-anonymous, l-diverse and t-close... ..so secure, right?

Caucas	787X X	HIV+	Flu
Asian/AfrAm	787X X	HIV-	Flu
Asian/AfrAm	787X X	HIV+	Shingles
Caucas	787X X	HIV-	Acne
Caucas	787X X	HIV-	Shingles
Caucas	787X X	HIV-	Acne

1.26 What Does Attacker Know?

Caucas	787X X	HIV+	Flu
Asian/AfrAm	787X X	HIV-	Flu
Asian/AfrAm	787X X	HIV+	Shingles
Caucas	787X X	HIV-	Acne
Caucas	787X X	HIV-	Shingles
Caucas	787X X	HIV-	Acne

1.27 k-Anonymity is Not Enough!

1.28 Syntactic Issues

- Syntactic
 - Focuses on data transformation, not on what can be learned from the anonymized dataset
 - “k-anonymous” dataset can leak sensitive information
- “Quasi-identifier” fallacy
 - Assumes a priori that attacker will not know certain information about their target
- Can increase levels of anonymity, but ...
 - Destroys utility of many real-world datasets